

STEM FELLOWSHIP 2025-2026 HIGH SCHOOL BIG DATA CHALLENGE

From Harvest to Food Security: Machine Learning Forecasts for the United States

Cathy Rong^{1, 2}, Safiya Taher¹, Sabrina Qian¹, Renisa Ghosh¹, and Arthur Jiang¹

¹Old Scona Academic, Edmonton, Alberta, Canada

²United World Colleges Red Cross Nordic, Flekke, Vestland, Norway

January 9, 2026

Corresponding author(s):

Cathy Rong (cathyrong11@gmail.com)

Word count (excluding figures, tables, references, and title page): 4455

Abstract

Food insecurity remains a persistent challenge, even in high-producing agricultural systems. Timely national-level indicators are needed to support monitoring and policy decisions. This study tests whether U.S. agricultural production signals can predict a composite Food Security Index (*FSI*) derived from food-system outcomes. We constructed an annual United States dataset spanning from 1970 to 2023 that integrates USDA NASS QuickStats measures for three major staple crops, corn, soybeans, and wheat (yield, production, area planted, and area harvested).

The prediction target is an *FSI* scaled from 0 to 1, where higher values indicate lower food insecurity, constructed through mathematical modeling of four FAOSTAT variables: dietary energy supply, protein supply, GDP per capita, and a food production index. We compared multiple supervised regression models using a fixed split of 1970 to 2012 data for training, and a testing period of 2013 to 2023. We further included a change-based formulation that predicts a year to year change to *FSI* levels (ΔFSI), with engineered shock and lag features.

A Random Forest model trained on the ΔFSI feature set performed the best, achieving a Root Mean Squared Error (RMSE) of about 0.066 on the 2013-2023 testing data and improving on a persistence baseline. Overall, results showed that the national crop statistics contain measurable predictive signals for annual food security dynamics and motive extensions.

Keywords:

Agriculture, Food Security Index,

1 Introduction

Even in highly productive agricultural systems, food insecurity is still a continuous challenge as national food security depends on more than the total output of crops [1]. Beyond aggregated supply, food security is a reflection on whether or not people can consistently access safe, sufficient, and nutritious food while the system remains stable under shocks [2]. As climate variability [3], market disruptions [4], and economic pressures interact, there is an increased interest in timely, interpretable national indicators that can be monitored over long periods in support of decision making. [5]

Agricultural production is one of the most consistently measured national variables. In the United States, corn, soybeans, and wheat are especially well documented as they dominate staple crop production and influence downstream

food and feed markets [6, 7]. Yearly changes in yield, acreage, and total production have the potential to capture biophysical constraints and economic land-use responses. However, directly linking production to national food security outcomes is not straightforward; food security also depends on economic access [8], supply chains [9, 10], and broader macro conditions, which introduce nonlinear relationships and time-lagged effects.

We created a composite food security index (*FSI*), computed from four different components derived from the Food and Agriculture Organization Corporate Statistical (FAOSTAT) Database [11]: dietary energy supply (kcal/capita/day), protein supply (g/capita/day), GDP per capita, and a food production index, to act as a summary measure score intended to track how “food-secure” the country is each year by using a series of macro indicators that are consistently available over longer time spans. It is a proxy for overall food security conditions rather than a measure for household hunger. Dietary energy supply provides information on food availability in the national supply [12]; protein supply shows a diet quality or adequacy dimension to the mix, considering beyond calories [13]; GDP per capita is used widely as a macro proxy for economic access and affordability where the higher the income, the higher the ability to purchase food, buffer price shocks, etc (generally) [14]; finally, the food production index is a broad domestic production availability signal by tracking the volume or value of edible, nutritious food crops relative to a base period [15].

This idea came from reviewing past composite food indices also built by selecting indicators and then normalizing and aggregating them. For example, the Global Food Security Index (GFSI) considers multiple indicators (food affordability, availability, quality and safety, sustainability and adaptation) over 113 different countries. [16, 17, 18, 19] These past indexes inspired us to create our *FSI* as a defensible, simple, and transparent measure, tailored reliably to what we are trying to solve.

2 Materials & Methods

2.1 Data and Study Design

We evaluated whether or not national staple crop production can predict a composite *FSI* for the United States (U.S.) using an annual time series (1970-2023) [20, 21]. Three different crop predictors were compiled from the United States Department of Agriculture (USDA) Na-




Figure 1: Study Flow Chart

tional Agricultural Statistics Service (NASS) Quickstats: corn, soybeans, and wheat, including yield, production, and harvested area for each crop ($N=1,944$) [11, 22]. The prediction target was an *FSI* scaled to $[0, 1]$, where a higher value signaled better food security.

In addition to crop exclusive statistics, we included a set of exogenous drivers to capture broader conditions that can influence food security and crop performance. U.S. food-price conditions were represented using a food Consumer Price Index (CPI) index and its derived annual change. The World Bank Commodity Price Data (“Pink Sheet”) indices (Energy, Agriculture, Food, Grains, Fertilizers, Metals & Minerals) and selected commodity price series (crude oil, maize, soybeans, and U.S. wheat prices: Hard Red Winter, HRW, and Soft Red Winter, SRW) acted as global input and commodity pressures.

All raw datasets were first restricted to the United States and then converted to annual values and merged by year to form a single master table. Non-numeric values were treated as missing and excluded only when required for the given model fit to maintain a consistent predictor-target alignment. The resulting dataset was then used for feature engineer-

ing (lags and year-over-year “shock” variables) and model evaluation under time-respecting train/test splits.

Figure 1 presents the flow chart for the machine learning model, including the data integration and model validation.

Notably, to minimize bias and maintain model integrity, no state data from QuickStats were included in the features. This approach ensured a clean separation of training, validation, and testing data to support robust model development and evaluation with the whole US context.

2.2 Data Preparation

We used the USDA NASS QuickStats API to extract our crop predictors [22]. Specifically, we sent requests that filtered records to U.S level, annual observations and targeted our three staple commodities (corn, soybeans, and wheat) as well as the required measures (yield, production, and area harvested). Each request was run through a Python code using the library ‘requests’ to call the API and saved each metric into its own structured dataset in CSV form. In addition, we downloaded exogenous series from the U.S. Bureau of Labor Statistics (CPI) [23], the World Bank Commodity Price Data (“Pink

Sheet”) [24, 25, 26] was used for indices and commodity prices and U.S. population totals from the World Bank population series.

Data cleaning and restructuring were also performed in Python [27], using the ‘pandas’ [28] and ‘NumPy’ [29] libraries for parsing, value conversion, reshaping, merging, and feature engineering (shock and lag value calculation). For consistency across all sources, all non-numeric values were dropped when required for a given model fit. In addition, we also used ‘pathlib’ [30] for file handling, and ‘scikit-learn’ utilities [31] to support model-ready preprocessing in our workflow. All of the cleaned series were then reshaped into a year-indexed table, at one row per year, and merged with the constructed FSI and exogenous predictors to create the final modeling dataset spanning from 1970 to 2023.

To ensure consistent scaling across variables with different units, all continuous features were normalized using the StandardScaler from scikit-learn’s preprocessing package. This normalization improved the comparability of features and optimized model performance. (detailed in Supplementary Table ??)

2.3 Targets and Feature Engineering

We created the FSI in multiple steps. For $FSI(t)$, we defined it as a single number between 0 and 1 for each year t , where the higher the value, the better food security (lower insecurity). We built it off of four annual national indicators:

1. $K(t)$: Dietary energy supply (kcal per capita per day)
2. $P(t)$: Protein supply (g per capita per day)
3. $G(t)$: GDP per capita (USD)
4. $F(t)$: Food production index

Because these four variables have different units and ranges, we first normalized each one to the same $[0, 1]$ scale as the final FSI over our full study cohort of 1970-2023, making them directly comparable. By using min-max normalization:

$$X_{\text{norm}}(t) = \frac{X(t) - \min(X)}{\max(X) - \min(X)}$$

for

$$X \in \{K, P, G, F\}$$

Each component lies in a range of $[0,1]$ after this.

Secondly, we combined all four components into one index by equal-weighted arithmetic mean. This is because we did not set any weights

to any of the components, making all of them of equal importance to one another.

$$FSI(t) = \frac{K_{\text{norm}}(t) + P_{\text{norm}}(t) + G_{\text{norm}}(t) + F_{\text{norm}}(t)}{4}$$

This, in turn, makes $FSI(t)$ a value within $[0,1]$, where a higher FSI value corresponds to stronger food security conditions, by our composite definition. At the same time, a drop in FSI would signify a deterioration driven by one or more components declining in relevance to their historical range.

While FSI provides an interpretable annual picture, its strong persistence over time reflects the idea that forecasting will be easier in year-to-year change increments rather than absolute levels. Therefore, we tested two prediction formulations: level and change-based. In the level formulation, models predict $FSI(t)$ directly, where $FSI(t)$ denotes the FSI in year t . In the change based formulation, models predict annual change $\Delta FSI(t) = FSI(t) - FSI(t-1)$ and reconstruct levels via $FSI'(t) = FSI(t-1) + \Delta FSI'(t)$. Where this year’s predicted FSI equals last year’s FSI plus the model’s predicted change this year. The FSI' means that value is a predicted FSI . Predictors included crop indicators (yield, production, and area harvested for corn, soybeans, and wheat) as well as exogenous variables in a climate and economic context.

To capture disruptions and delayed effects, we manually created shock features and lags for all time-varying predictors. Shocks were computed on a year by year basis as changes .(percentage changes for strictly positive variables) We also included 1 to 2 year lags of both raw and shocked level terms. This design allows our model to learn immediate versus delayed associations between production, external conditions, and changes in national food security.

2.4 Model Selection

To accurately test whether or not all of the predictors could contain a usable signal for food security, we employed multiple supervised machine learning regression models to predict for our FSI . The model had a time-respective training and testing data split where 1970 to 2012 data was used for training and 2013 to 2023 data was used for testing. The model was benchmarked against a persistent baseline that forecasts the next year’s value.

We compared a small set of regression models to test, each with different strengths. These included a regularized Ridge regression linear model as well as a non linear Random Forest

regression ensemble model. The Ridge model provides a transparent benchmark that handles all correlated predictors through regularization; the Random Forest model can help capture non-linear interactions between variables without requiring a separate fixed functional form. To evaluate which model allowed for easier forecasting, our two formulation processes were used, and the best-performing approach was selected based on predictive accuracy and the held-on test period.

2.5 Performance Metrics

Predictive accuracy on held-out years was measured using Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), and R^2 values, reported in relevance to the persistence baseline. For the second formulation, the change-based formulation, we assessed directional performance additionally by classifying whether ΔFSI was bigger than 0 and then computing a confusion matrix and balanced accuracy along with it. Receiver Operating Characteristic Area Under the Curve (ROC AUC) and average precision were also calculated using the changed, predicted scores.

3 Results

Through national staple-crop statistics and exogenous economic indicators, we evaluated the predictive perforation of the proposed Ridge and Random Forest models for estimating food security in the United States. In this section, we see how well our models predict U.S. food security in the held-out period and what signals drive said predictions.

According to the results across the tested models, the Random Forest model, trained on ΔFSI , was selected as the primary approach as it provided the most reliable out-of-sample tracking of FSI relative to the alternative models and our persistence baseline. Therefore, we present the results from the Random Forest model first, followed by diagnostics and an auxiliary directional evaluation.

Figure 2 shows a comparison between the observed FSI in the 2013 to 2023 test period and the predictions from the Random Forest model trained on annual changes (ΔFSI). In addition, a persistence baseline that simply carries forward the previous FSI ($FSI(t-1)$) is also compared. We set the y-axis to start at 0.60, rather than 0, to focus on the 2013-2023 range and make small but meaningful differences between the observed FSI , model predictions, and the persistence baseline easier to see. Overall,

the Random Forest model is able to forecast patterns following the same broad trajectory as the observed series, including a mid-decade decline and substantial rise right after 2017. This suggests that the model captures meaningful year by year movements rather than only the trend in the long run. The persistence baseline also helps track the general direction; however, it responds more slowly to turning points because it assumes that there has been no change beyond last year’s level. Differences between predicted and observed curves are most visible around the sharper transitions, where the model either slightly over or underestimates the magnitude of change. But, the close alignment across the vast majority of the years shows that crop and “shock and lag” features contain usable signals for national food-security dynamics.

Figure 3 plots the same test-year values as Figure 2, but collapses them into point pairs where each dot is one year with the x-axis being the observed FSI and the y-axis as the predicted FSI . Points closer to the line of best fit means better agreement. In our test set, our predictions are aligned with the observations with a correlation value of around 0.82, with 6 out of 11 years falling within 0.05 of the true value and 4 out of 11 within 0.02. Similarly to Figure 1, the largest mismatches occur in the volatile years: 2017 being one as it is underpredicted.

Figure 2 is a time-series story which shows whether the model follows a yearly trajectory with turning points compared to a persistence baseline. On the other hand, Figure 3 is a calibration, or fit, check as it ignores ordering in time and visualizes how close predictions in our model are to the correct values overall.

Figure 4 shows the residual of the model, an annual year-by-year prediction error in the testing period. This is defined by the equation: $r(t) = FSI_{\text{actual}}(t) - FSI_{\text{pred}}(t)$. Positive results indicate underprediction, where the observed FSI was higher than the estimation made by the model, while negative residuals indicate overprediction. The centre line at Residual = 0 means that there is “no error” where the model predicted the FSI exactly for that year. Across 2013 to 2023, residuals are generally small and only fluctuate around zero, where the mean residual is about +0.016 and the median is around -0.016. This suggests that the model is not consistently biased in one signal direction over the entire testing window.

However, the plot also shows the years that drive most of the error. The largest underprediction occurs in 2017, where the residual is +0.150, with an actual FSI value of 0.738 and a predicted value of 0.588. The largest overpredic-




Figure 2: Observed vs Predicted ΔFSI Over Time: Random Forest vs Persistence (2013-2023)




Figure 3: Predicted vs Observed FSI in the Test Period (2013-2023)




Figure 4: Residuals Over Time (Observed - Predicted *FSI*), 2013-2023

tion happens in 2016, with a residual of -0.108 and a split of 0.599 as the actual values versus the predicted 0.707. These spikes cluster in the sharp transitions in the *FSI* trajectory, further showing that the model struggles to adapt when food security shifts abruptly. Overall, the residual pattern supports the interpretation that our model is able to track gradual movement well but extreme changes still remain the biggest hardships.

The importance of features for predicting *FSI* was evaluated based on the model coefficients. Features with larger absolute coefficient values were deemed more significant, indicating a stronger influence on the target outcome.

Figure 5 shows which predictors the Random Forest model relied on the most when forecasting ΔFSI . This importance ranking is dominated by production-side crop variables and short-term temporal structure, especially on the wheat and corn variants. The most influential predictor was “wheat harvested area (t-1)” with an importance of around 0.064, followed closely by “corn yield (level)” at around 0.060. Other high ranking features are also lagged crop quantities such as “wheat product (t-2)” (≈ 0.040), “wheat yield (t-2)” (≈ 0.036), and “wheat production (t-1)” (≈ 0.030). This indicates that the

model draws heavily on one to two year delayed information rather than contemporaneous levels. This means that lag features (labels such as t-1 and t-2) are simply predictor values that are taken from previous years, allowing the model to capture delayed effects where impacts on food security unfold over time.

In addition, shock-style variables also appear repeatedly. Features labeled with “(YoY %)” are a representation of year by year percent change, meaning how much a variable has increased or decreased compared to the previous year. In Figure 4, these YoY % are the shock variables to help the model distinguish between short-run departures and typical conditions. The appearance of “soybeans yield (YoY %)” (≈ 0.033), “soybeans production (YoY %)” (≈ 0.029), and “corn yield (YoY %)” (≈ 0.026) among the top predictors suggests that year by year sudden changes do contain useful signal for changes in national food security.

Finally, market price factors also contribute but do not dominate: “U.S. wheat price (SRW) (YoY %)” (≈ 0.024), “maize price (YoY %)” (≈ 0.019), and “U.S. wheat price (HRW)” (YoY %) (≈ 0.016) are included in the top 15, which implies that these commodity-price shocks help provide an additional contest beyond production




Figure 5: Top 15 Feature Importances: Random Forest ΔFSI

volumes. Most importantly, these reflect predictive contributions rather than just casual effects. Figure 5 indicates which variables were most useful in reducing prediction error when the Random Forest learned patterns ΔFSI .

Table 1 shows how sensitive the Random Forest ΔFSI model is to the way we encoded short-run disruptions and delayed effects in our predictors. We varied two feature types: our shock features (year by year percentage changes, “_yoy_pct”) and our lag features (prior year values, “_lag1” or “_lag2”). This setup helps isolate whether the model benefits more from immediate year by year changes, historical memory, or both.

The pattern is evident: accuracy suffers when temporal structure is removed, but accuracy is consistently improved when shocks are combined with short lags. Error increases significantly (RMSE 0.0788, R^2 0.5478) when both shocks and lags are eliminated (“no shocks + no lags”), and the use of only shocks or only lags stays weak (RMSE 0.0791 and 0.0809, respectively). On the other hand, when shocks and lags are combined, error is significantly reduced and performance approaches the persistence benchmark (RMSE 0.0699). Shocks + lag1+lag2 (no lagged shocks) is the grid’s best-performing configura-

tion (RMSE 0.0701, MAE 0.0509, R^2 0.6419), and lag2 shock setups also enhance directional performance (balanced accuracy 0.7143). Overall, Table 1 demonstrates that the model best captures the dynamics of national food security when it can learn from both YoY shocks and 1-2 year delayed effects rather than just static levels alone.

Figure 6 is an evaluation on whether the model is able to predict the direction of food security change on an annual scale by framing the outcome as a binary label. Our definitions included: Positive if $\Delta FSI > 0$, if FSI improved from the previous year, and Negative if $\Delta FSI \leq 0$. This figure includes two visualizations. The ROC curve (6a) helps summarize how well the model is able to distinguish between these two classes across all possible decision thresholds. With an AUC of 0.68, it indicated modest to moderate discrimination, which is better than chance (0.50), but also not the most strong. Our second visualization (6b) is the confusion matrix which is shown as percentages with the predicted classes (normalized by column), so each column is a reflection of how accurate the model’s predicted label is. When the model predicts a Negative year, 75.0% of those predictions are truly negative (TN) and 25.0% are actually positive

Table 1: Shock-Lag Sensitivity Analysis for the Random Forest ΔFSI Model (Test: 2013-2023)

Configuration	Shocks	Lag Depth	Lagged Shocks	RMSE	MAE	R ²	Balanced Accuracy
Persistence baseline (FSI(t)=FSI(t-1))	—	—	—	0.0699	0.0554	0.6443	
Shocks + lag1+lag2 (no lagged shocks)	On	2	Off	0.0701	0.0509	0.6419	0.7143
Shocks + lag1+lag2	On	2	On	0.0704	0.0497	0.6393	0.7143
Shocks + lag1	On	1	On	0.0712	0.0517	0.6308	0.6429
No shocks + no lags	Off	0	Off	0.0788	0.0646	0.5478	0.5
Shocks only	On	0	Off	0.0791	0.0592	0.5438	0.6429
Lag1 only	Off	1	Off	0.0809	0.071	0.5227	0.5




Figure 6: ROC Curve and Confusion Matrix for Predicting $\Delta FSI > 0$

(FN). When the model predicts a Positive year, 85.7% are truly positive (TP) and 14.3% are false positives (FP).

Table 2 is a comparison of directional performance across a full set of models by evaluating whether each approach correctly predicts the sign of annual food security change, where Positive is $\Delta FSI > 0$ and Negative is ≤ 0 . While the earlier results focus exclusively on the Random Forest ΔFSI model for detailed diagnostics, Table 2 is a showcase of the other models we tried to test with by contextualizing how alternative methods and feature sets perform on the same 2013 to 2023 test window.

The strongest directional performance was achieved by the Random Forest model using our fully engineered feature set (ΔFSI : crops + macro + shocks + lag), which achieved a balanced accuracy of 0.8036, with sensitivity 0.8571 and specificity 0.7500. With ΔFSI : macro + shocks + lags, the Random Forest model achieved a balanced accuracy of 0.7321, with sensitivity 0.7143 and specificity 0.7500. For

ΔFSI : crops + shocks + lags, Gradient Boosting reaches a balanced accuracy of 0.6786 (sensitivity 0.8571, specificity 0.5000), suggesting it is comparatively good at catching improvement years but less consistent at correctly identifying non-improvement years. Under the same reduced data set, RidgeCV yields a balanced accuracy of 0.6250 with perfect sensitivity (1.0000) but low specificity (0.2500).

4 Discussion

Overall, our results showed that national crop yield, production, and harvest area statistics helps signal tracking for U.S. food security over time. The Random Forest ΔFSI model consistently provided the most reliable results out of all the tested approaches. We discovered that supervised models can learn relationships between staple-crop indicators and a composite FSI using an annual U.S.-only dataset (1970 to 2023), particularly when the target is modeled

Table 2: Directional Model Comparison for ΔFSI :
Balanced Accuracy, Sensitivity, and Specificity (Test: 2013-2023)

Setup	Model	Balanced Accuracy	Sensitivity	Specificity
ΔFSI : crops+macro+shocks+lags	RandomForest	0.8036	0.8571	0.75
ΔFSI : macro+shocks+lags	RandomForest	0.7321	0.7143	0.75
ΔFSI : crops+shocks+lags	GradientBoosting	0.6786	0.8571	0.50
ΔFSI : crops+shocks+lags	RidgeCV	0.625	1	0.25
Baseline	Persistence ($\Delta = 0$)	0.5	0	1

as year-to-year change (ΔFSI) rather than absolute levels. The Random Forest ΔFSI framework performed competitively against a strong persistence benchmark and generally followed the observed trajectory during the 2013 to 2023 holdout period. This lends credence to the notion that national crop conditions frequently correspond with more general food-system outcomes, making them useful for forecasting and monitoring.

A major takeaway is that our shock and lag feature engineering mattered as much as model choice. The shock-lag sensitivity results show that removing both YoY shock terms and lag structure noticeably worsened performance, whereas including both consistently improved results. This is plausible in food systems where production outcomes, input costs, and price conditions influence food availability and stability with delays rather than instantaneously. In addition, lagged crop quantities (“t-1” and “t-2”) and shock terms (“YoY %”) were amongst some of the strongest predictors in Figure 4, implying that the model benefits from both the “memory” of prior conditions as well as a detection of sudden year by year disruptions.

At the same time, our results clarify what exclusively crop-based prediction can and cannot do. The residual analysis shows that the largest errors occur around years with sharper transitions, and the directional evaluation shows that for up or down changes in ΔFSI , there is only moderate discrimination with an AUC of 0.68. This is consistent with the broader reality of food security as it is not solely determined by production; affordability, distribution, macroeconomic shocks, and policy/trade dynamics all have potential to shift outcomes even if harvest data shows strength. As a result, production and commodity indicators are most useful for identifying medium-term trends and gradual movement, and less trustworthy when changes in access or policy predominate in the food security signal.

Since directional classification is more diffi-

cult than level prediction in this context, the AUC is not higher. First, a few borderline years can significantly alter the ROC curve and lower AUC because the test window only spans 11 years. Second, the positive/negative boundary is susceptible to small noise in either the predictors or the constructed FSI itself because many year-to-year changes in a national index are small (near $\Delta FSI=0$). Third, crop and commodity indicators may track the overall movement of FSI without consistently capturing the precise up/down. This is because, even when agricultural production contains a signal, the sign of food-security change yearly also depends on drivers that are only partially represented in our features, such as policy shifts, trade dynamics, household prices, distribution, and access constraints. For these reasons, our ROC/AUC results are a supporting diagnostic where the model shows some ability to detect improvement years.

However, when interpreting these results, a number of limitations should be taken into account. Due to the short test window from 2013 to 2023, performance metrics are sensitive to a few borderline and inconsistent years. Localized shocks that do not significantly register in aggregates may be missed by the national annual design, which smoothes within-country variation. Although the FSI offers a consistent 0 to 1 target, any index depends on design decisions that may impact year-to-year dynamics. The FSI itself is a constructed composite based on modeled transformations of FAOSTAT variables. Lastly, despite the fact that we included a number of macro and commodity-price series, the feature set leaves out important factors that affect food security, such as measures of household income and poverty, changes in policy and assistance, trade volumes, inventories and stock-to-use conditions, and distribution restrictions.

Otherwise, our findings have a direct relevance to policy. Firstly, it shows the use of routinely and easily accessible crop statics as a part of an early-warning style monitoring. Our approach

is best-suited to flagging when the U.S. system is deviating from recent norms as we use the best-forming setups including YoY shocks and short lags. This can help agencies and decision-makers prioritize attention, scenario planning, and preparedness. Secondly, our limitations are also equally informative. Moderate ROC/AUC performance and outlier years, like in 2017, reinforce the idea that supply-exclusive improvements alone does not guarantee better food security. Implying that yield-boosting or acreage-expansion policies should be paired with measures that target affordability and access, such as food price stabilization, safety net strengthening, or improving distribution reliability, our limitations actually help and inform those in power of the main goal of translating production strength into real gains in food-security.

In conclusion, crop indicators by themselves cannot fully explain food security, a multi-layer commodity, but they are sufficiently informative to support tracking and monitoring at the national level, particularly when modelled with shocks and lags. The most compelling evidence from our work is that expanding the predictor set toward access, trade, and policy drivers is the most promising route to policy-relevant forecasting, and that year by year changes in food security are more predictable when models can learn from both short-run disruptions and delayed effects.

5 Conclusion

Even in a high-producing agricultural system, such as the United States, national food security can also shift, which makes timely monitoring and early warnings especially important. In conclusion, crop indicators by themselves cannot fully explain food security, but they are sufficiently informative to support tracking and monitoring at the national level, particularly when modelled with shocks and lags. The most compelling evidence from this work is that expanding the predictor set toward access, trade, and policy drivers is the most promising route to policy-relevant forecasting, and that year to year changes in food security are more predictable when models can learn from both short-run disruptions and delayed effects. Future work should focus on further increasing our predictor set to include more variables such as access, trade, and policy drivers and move toward uncertainty-aware forecasts. That allows for decision-makers to more effectively allocate resources and respond earlier to conditions that threaten national food security. This approach has the potential to respond to noticed issues faster and

more effectively identifying food insecurity risk.

Acknowledgements

The team would like to thank the support from Dr. Jianshan Chen in University of Alberta. We would also like to thank STEM Fellowship for providing us the opportunity to participate in the HSBDC.

References

- [1] Michael A. Long, Lara Gonçalves, Paul B. Stretesky, and Margaret Anne Defeyer. Food insecurity in advanced capitalist nations: A review. *Sustainability*, 12(9), 2020.
- [2] Aikaterini Kavallari, Thomas Fellmann, and S. Hubertus Gay. Shocks in economic growth = shocking effects for food security? *Food Security*, 6(4):567–583, August 2014.
- [3] Robert H. Beach, Allison M. Thomson, and Bruce A. McCarl. Climate change impacts on us agriculture. *IATRC Public Trade Policy Research and Analysis Symposium*, 2010. Available from: <https://ageconsearch.umn.edu/record/91393>.
- [4] Regis Barnichon, Christian Matthes, and Alexander Ziegenbein. Are the effects of financial market disruptions big or small? *The Review of Economics and Statistics*, 104(3):557–570, 05 2022.
- [5] J. L. Hatfield, K. J. Boote, B. A. Kimball, L. H. Ziska, R. C. Izaurralde, D. Ort, A. M. Thomson, and D. Wolfe. Climate impacts on agriculture: Implications for crop production. *Agronomy Journal*, 103(2):351–370, 2011.
- [6] Jasmien De Winne and Gert Peersman. Macroeconomic effects of disruptions in global food commodity markets: Evidence for the united states. *Brookings Papers on Economic Activity*, 2016(2):183–286, 2016. Fall 2016.
- [7] Maayan Kreitzman, Eric Toensmeier, Kai M. A. Chan, Sean Smukler, and Navin Ramankutty. Perennial staple crops: Yields, distribution, and nutrition in the global food system. *Frontiers in Sustainable Food Systems*, Volume 4 - 2020, 2020.
- [8] N. M. A. Manap and N. W. Ismail. Food security and economic growth. *International Journal of Modern Trends in Social Sciences*, 2(8):108–118, 2019.

- [9] Noé J. Nava, William Ridley, and Sandy Dall’erba. A model of the u.s. food system: What are the determinants of the state vulnerabilities to production shocks and supply chain disruptions? *Agricultural Economics*, 54(1):95–109, 2023.
- [10] B. Gail Smith. Developing sustainable food supply chains. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 363(1492):849–861, 08 2007.
- [11] FAOSTAT. Country indicators, 2025. Available from: <https://www.fao.org/faostat/en/#country>.
- [12] M Nubé. Confronting dietary energy supply with anthropometry in the assessment of undernutrition prevalence at the level of countries. *World Development*, 29(7):1275–1289, 2001.
- [13] Maeve Henchion, Maria Hayes, Anne Maria Mullen, Mark Fenelon, and Brijesh Tiwari. Future protein supply and demand: Strategies and factors influencing a sustainable equilibrium. *Foods*, 6(7), 2017.
- [14] Peter Warr. Food insecurity and its determinants. *Australian Journal of Agricultural and Resource Economics*, 58(4):519–537, 2014.
- [15] İbrahim Bozkurt and MUHAMMED VEYSEL Kaya. Agricultural production index: International comparison. *Agricultural Economics (Czech Republic)*, pages 236–245, 2021.
- [16] Economist Impact Expert panel. Global food security index 2022: Exploring challenges and developing solutions for food security across 113 countries, 2022. Available from: <https://impact.economist.com/sustainability/project/food-security-index/>.
- [17] Farhana Bibi and Azizur Rahman. An overview of climate change impacts on agriculture and their mitigation strategies. *Agriculture*, 13(8), 2023.
- [18] Rosana Salles-Costa, Aline Alves Ferreira, Ruben Araujo de Mattos, Michael E. Reichenheim, Rafael Pérez-Escamilla, Juliana de Bem-Lignani, and Ana Maria Segall-Corrêa. National trends and disparities in severe food insecurity in brazil between 2004 and 2018. *Current Developments in Nutrition*, 6(4), Apr 2022.
- [19] Akinniyi Akinwande. Nigeria agri-econ indicators (1950-2023), 2024. Available from: <https://www.kaggle.com/datasets/akinniyiakinwande/nigeria-agri-econ-indicators-1950-2023>.
- [20] Justin Oh. Total emissions per country (2000-2020), 2023. Available from: <https://www.kaggle.com/dsv/5097715>.
- [21] Hannah Ritchie, Pablo Rosado, and Max Roser. Agricultural production. *Our World in Data*, 2023. Available from: <https://ourworldindata.org/agricultural-production>.
- [22] U.S. Department of Agriculture, National Agricultural Statistics Service. Quick Stats Agricultural Database. <https://quickstats.nass.usda.gov/>, 2025. Accessed via the Quick Stats API: <https://quickstats.nass.usda.gov/api>.
- [23] Varshita Nalluri. Crop price prediction dataset, 2024. Available from: <https://www.kaggle.com/datasets/varshitanalluri/crop-price-prediction-dataset>.
- [24] World Bank Group and partners. Environment - databank, 2021. Available from: <https://data.worldbank.org/topic/environment>.
- [25] World Bank Group and partners. World development indicators - databank, 2015. Available from: <https://databank.worldbank.org/id/db4738db>.
- [26] World Bank Group and partners. World bank solutions to food insecurity - food security update, 2025. Available from: <https://www.worldbank.org/en/topic/agriculture/brief/food-security-update>.
- [27] Python Software Foundation. *Python Language Reference, version 3.11*, Oct 2024. <https://docs.python.org/3/reference/index.html>.
- [28] The pandas development team. pandas-dev/pandas: Pandas, February 2020.
- [29] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, and Ralf Gommers et. al. Array programming with NumPy. *Nature*, 585(7825):357–362, September 2020.
- [30] Python Software Foundation. The Python™ Programming Language. Documentation 3.11, 2026. Specifically for the `pathlib` module.

- [31] F. Pedregosa, Gaël Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in python. *Journal of machine learning research*, 12(Oct):2825–2830, 2011.